

title: Docker Monitoring Integration description: Container performance monitoring - metrics, event, log collection and parsing for Docker

The Sematext Integration for Docker uses Sematext Agent and the open-source Logagent for the collection of container logs.

We at Sematext aim to save you time and effort by giving you a strong starting point for monitoring Docker. You will **not** have to:

- figure out which metrics to collect and which ones to ignore
- give metrics meaningful labels
- hunt for metric descriptions in the docs so that you know what each of them actually shows
- build charts to group metrics that you really want on the same charts, not N separate charts
- figure out, for each metric, which aggregation to use (min? max? avg? something else?)
- build dashboards to combine charts with metrics you typically want to see together
- set up basic alert rules

We set it up for you, out-of-the-box!

## Docker Monitoring with Sematext Agent

Sematext Agent collects various metrics about hosts and ships that to Sematext Cloud.

First create a Docker Monitoring App in Sematext Cloud.

### Docker

You install the Agent simply by running one Docker command. This will start the Agent as a Docker container on your host.

```
docker run -d --restart always --privileged -P --name st-agent \
-v /:/hostfs:ro \
-v /sys/kernel/debug:/sys/kernel/debug \
-v /var/run:/var/run/ \
-v /proc:/host/proc:ro \
-v /etc:/host/etc:ro \
-v /sys:/host/sys:ro \
-v /usr/lib:/host/usr/lib:ro \
-e CONTAINER_TOKEN=84fbc37e-a0fb-418c-9bcb-ea7c763dd9ac \
-e INFRA_TOKEN=6377be8e-8441-46de-85dc-11ee3646c3de \
-e REGION=US \
-e JOURNAL_DIR=/var/run/st-agent \
-e LOGGING_WRITE_EVENTS=false \
-e LOGGING_REQUEST_TRACKING=false \
```

```
-e LOGGING_LEVEL=info \
-e NODE_NAME=`hostname` \
-e CONTAINER_SKIP_BY_IMAGE=sematext \
sematext/agent:latest
```

## Docker Compose

If you prefer adding the Agent in a docker-compose configuration, here's how you do it.

```
# docker-compose.yml
version: '3'
services:
  sematext-agent:
    image: 'sematext/agent:latest'
    environment:
      - affinity:container!=*sematext-agent*
      - CONTAINER_TOKEN=84fbc37e-a0fb-418c-9bcb-ea7c763dd9ac
      - INFRA_TOKEN=6377be8e-8441-46de-85dc-11ee3646c3de
      - REGION=US
      - JOURNAL_DIR=/var/run/st-agent
      - LOGGING_WRITE_EVENTS=false
      - LOGGING_REQUEST_TRACKING=false
      - LOGGING_LEVEL=info
      - NODE_NAME=$HOSTNAME
      - CONTAINER_SKIP_BY_IMAGE=sematext
    cap_add:
      - SYS_ADMIN
    restart: always
    volumes:
      - '/:/hostfs:ro'
      - '/var/run:/var/run/'
      - '/sys/kernel/debug:/sys/kernel/debug'
      - '/proc:/host/proc:ro'
      - '/etc:/host/etc:ro'
      - '/sys:/host/sys:ro'
      - '/usr/lib:/host/usr/lib:ro'
```

Then you can run one command to start the Agent.

```
docker-compose up -d
```

## Docker Swarm

If you're running a Docker Swarm cluster, it's just as easy to run a Docker Swarm service.

```

docker service create --mode global --name st-agent \
  --restart-condition any \
  --mount type=bind,src=/,dst=/hostfs,readonly \
  --mount type=bind,src=/var/run,dst=/var/run/ \
  --mount type=bind,src=/usr/lib,dst=/host/usr/lib \
  --mount type=bind,src=/sys/kernel/debug,dst=/sys/kernel/debug \
  --mount type=bind,src=/proc,dst=/host/proc,readonly \
  --mount type=bind,src=/etc,dst=/host/etc,readonly \
  --mount type=bind,src=/sys,dst=/host/sys,readonly \
  -e NODE_NAME={{.Node.Hostname}} \
  -e CONTAINER_TOKEN=84fbc37e-a0fb-418c-9bcb-ea7c763dd9ac \
  -e INFRA_TOKEN=6377be8e-8441-46de-85dc-11ee3646c3de \
  -e REGION=US \
  -e JOURNAL_DIR=/var/run/st-agent \
  -e LOGGING_REQUEST_TRACKING=false \
  -e LOGGING_WRITE_EVENTS=false \
  -e LOGGING_LEVEL=info \
  -e PKG_ENABLED=false \
  sematext/agent:latest

```

If you like using docker stack, editing the docker-compose.yml from above slightly you'll have a working configuration.

```

version: "3"
services:
  agent:
    image: sematext/agent:latest
    deploy:
      mode: global
      labels: [APP=AGENT]
      restart_policy:
        condition: any
        delay: 1s
    cap_add:
      - SYS_ADMIN
    restart: always
    environment:
      - affinity:container!=*sematext-agent*
      - CONTAINER_TOKEN=84fbc37e-a0fb-418c-9bcb-ea7c763dd9ac
      - INFRA_TOKEN=6377be8e-8441-46de-85dc-11ee3646c3de
      - JOURNAL_DIR=/var/run/st-agent
      - LOGGING_WRITE_EVENTS=false
      - LOGGING_REQUEST_TRACKING=false
      - LOGGING_LEVEL=info
      - NODE_NAME=$HOSTNAME
      - CONTAINER_SKIP_BY_IMAGE=sematext
      - REGION=US

```

```

    - PKG_ENABLED=false
  volumes:
    - "/:/hostfs:ro"
    - "/var/run:/var/run/"
    - "/usr/lib:/host/usr/lib"
    - "/sys/kernel/debug:/sys/kernel/debug"
    - "/proc:/host/proc:ro"
    - "/etc:/host/etc:ro"
    - "/sys:/host/sys:ro"

```

Then you run:

```
docker stack deploy -c docker-compose.yml <name>
```

The Sematext Agent will start collecting dozens of key metrics right away, and start showing you the performance and health of your Docker containers immediately.

## Collected Docker Metrics

The Sematext Agent will collect the following container and host metrics.

### Host Metrics

- CPU
- memory
- disk
- network
- processes
- containers
- orchestrator platforms

### eBPF Support

To gain deep **insight into the Linux kernel**, Sematext Agent relies on **eBPF** to implant **instrumentation points**, which means to **attach eBPF programs to kprobes** on kernel functions. This ensures a very efficient and powerful system exploration approach with better network tracing and negligible overhead.

### Service Auto-Discovery

Sematext Agent can **auto-discover services** deployed on physical/virtual hosts and containers. It also collects data about your infrastructure to provide you with infrastructure inventory reports. It collects events from different sources such as OOM notifications, container or Kubernetes events.

## Container Metrics

- Container runtime agnostic discovery and monitoring
  - Containers are discovered from cgroups hierarchies
  - Supports Docker and Rkt container engines
- Container metrics fetched directly from cgroups
  - CPU usage
  - Disk space usage and IO stats
  - Memory usage, memory limits, and memory fail counters
  - Network IO stats
- Collection of host inventory information
  - Host kernel version/system information
  - Information about installed software packages
- Collection of container metadata
  - Container name
  - Image name
  - Container networks
  - Container volumes
  - Container environment
  - Container labels including relevant information about orchestration
  - Kubernetes metadata such as Pod name, UUID, Namespace
  - Docker Swarm metadata such as Service name, Swarm Task etc.
- Collection of container events
- Docker events such as start/stop/die/volume mount, etc.
- Kubernetes events such as Pod status changes deployed, destroyed etc.
- Tracking deployment status and Pod restarts over time

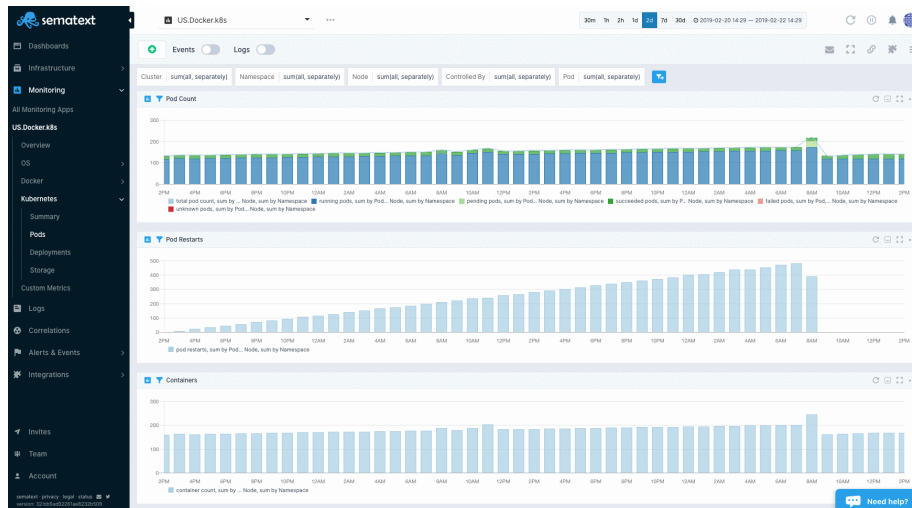
That is a lot of information and **Sematext organizes this information in reports for infrastructure monitoring, container monitoring, and Kubernetes cluster monitoring.**

## Docker Alerting

To save you time Sematext automatically creates a set of default alert rules such as alerts for low disk space. You can create additional alerts on any metric.

There are 3 types of alerts in Sematext:

- **Heartbeat alerts**, which notify you when a server is down
- **Threshold-based alerts** that notify you when a metric value crosses a predefined threshold
- **Alerts** based on statistical **anomaly detection** that notify you when metric values suddenly change and deviate from the baseline



## Docker Events

Events reflect changes in your infrastructure, from node restarts to container deployments, or changes in running containers. Events can track every Docker command. **Sematext Agent collects Events from the Docker Engine and Kubernetes API.** Whenever something goes wrong in your container stack, you can **correlate Logs or Metrics with the time of Docker events!**

Here's the list of Docker container events Sematext collects:

### Container lifecycle events

- Create – when a container is created
- Start – when a container starts
- Restart – when a container gets restarted
- Stop – when a container stops
- Oom – when a container runs out of memory
- Pause – when a container gets paused
- Unpause – when a container continues to run after a pause
- Die – when the main process in a container dies
- Kill – when the container gets killed
- Destroy – when a container gets destroyed

### Container runtime events

- Commit – when changes to the container filesystem are committed. Modifying deployed containers in production is not a common practice, therefore the commit could indicate a “hack” and should be watched carefully.
- Copy – when files are copied from/to a container. Could indicate a potential data leak.

- Attach – when a process connects to container console – somebody is reading your container logs
- Detach – when a process disconnects from container console streams
- Exec – when a command is executed in container console, very helpful to investigate in potential hacker attacks
- Export – when a container gets exported
- Health\_status – when health\_status is checked
- Rename – when a container gets renamed
- Resize – when a container gets resized
- Top – when somebody list top processes in a container
- Update – when a container is updated e.g. with new labels

### **Container image events**

- Delete – when an image gets deleted
- Import – when an image gets imported
- Load – when an image is loaded
- Pull – when an image is pulled from a registry
- Push – when an image is pushed to a registry
- Save – when an image is saved
- Tag – when an image is tagged with labels
- Untag – when an image tag is removed

### **Container plugin events**

- Enable – when a plugin gets enabled
- Disable – when a plugin gets disabled
- Install – when a plugin gets installed
- Remove – when a plugin gets removed

### **Container volume events**

- Create – when a volume is created
- Destroy – when a volume gets destroyed
- Mount – when a volume is mounted to a container
- Unmount – when a volume is removed from a container

### **Container network events**

- Create – when a network is created
- Connect – when a container connects to a network
- Remove – when the network is removed
- Destroy – when a network is destroyed
- Disconnect – when a container disconnects from a network

## Docker daemon events

- Reload

## Docker services, nodes, secrets, and config events

- Create – on the creation of a resource
- Remove – on the removal of a resource
- Update – on the creation of a resource

## Metrics Overview

The following information is collected and transmitted to Sematext Cloud or Sematext Enterprise.

Type

Description

Operating System Metrics

Host machine metrics

CPU Usage

Memory Usage

Network Stats

Disk I/O Stats

Container Metrics/Stats

CPU Usage / limits

Memory Usage / Limits / Fail Counters

Network Stats

Disk I/O Stats

Events

Agent Startup Event

server-info – created by spm-agent framework with node.js and OS version info on startup. Please note the agent is implemented in node.js.

Docker-info – Docker Version, API Version, Kernel Version on startup

Docker Events

Container Lifecycle Events| create, exec\_create, destroy, export, ...



Container Runtime Events

die, exec\_start, kill, pause, restart, start, stop, unpause, ...

Docker Logs

Default Fields

hostname / IP address

container id

container name

image name

message

Log formats

(detection and log parsers)

NGINX

APACHE httpd, Kafka, Solr, HBase, Zookeeper, Cassandra

MySQL

MongoDB

Redis

Elasticsearch

NSQ / Nsq.io

patterns are maintained here:

<https://github.com/sematext/logagent-js>

JSON, Plain Text

### **Supported Platforms**

- Docker Engine  $\geq 17.0.0$
- Platforms using Docker:
  - Docker Cloud
  - Docker Data Center
  - Kubernetes
  - Mesos
  - CoreOS
  - Rancher
  - Amazon ECS
  - Red Hat OpenShift

## Metrics Fields

Name	Type	Unit	Numeric Type	Label	Description
containermemory.usage	gauge	bytes	long	memory	container memory usage in bytes
containermemory.fail.count	counter		long	memory	the number of times that memory cgroup limit was exceeded
containermemory.limit	gauge	bytes	long	memory	the max allowed memory limit for the container cgroup
containermemory.limit.soft	gauge	bytes	long	soft memory limit	soft memory limit represents the initial memory reservation for the container
containermemory.rss	gauge	bytes	long	RSS memory	number of bytes of anonymous (file unmapped memory) and swap cache memory
containermemory.page.usage	gauge	bytes	long	cache memory	number of bytes of page cache memory
containermemory.pages.in	counter		long	memory pages in	memory pages in,description=the number of events each time the page is accounted to the cgroup
containermemory.pages.out	counter		long	memory pages out	memory pages out,description=the number of events each time a page is unaccounted from the cgroup
containermemory.pages.faults	counter		long	memory page faults	the number of page faults accounted to the cgroup

Name	Type	Unit	Numeric Type	Label	Description
containermemory.major	memory	pages	long	major mem-ory page faults	the number of major page faults accounted to the cgroup
containermemory.swap	memory	bytes	long	swap	the number of bytes of swap usage
containermemory.swap.limit	memory	bytes	long	swap limit	the swap memory usage limit
containerio.read	io	bytes	long	disk read	the number of bytes read from the disk
containerio.read.time	io	time	long	disk read time	the total amount of time (in nanoseconds) between request dispatch and request completion
containerio.read.wait	io	time	long	disk read wait time	total amount of time the IO operations for this cgroup spent waiting in the scheduler queues
containerio.write	io	bytes	long	disk write	the number of bytes written to the disk
containerio.write.time	io	time	long	disk write time	the total amount of time (in nanoseconds) between request dispatch and request completion
containerio.write.wait	io	time	long	disk write wait time	total amount of time the IO operations for this cgroup spent waiting in the scheduler queues

Name	Type	Unit	Numeric Type	Label	Description
containerengine.weights	long		long	disk io weight	specifies the relative proportion of block I/O access ranging from 100 to 1000
containerengine.cpu.percent	double	%	double	CPU usage	container CPU usage
containerengine.cpu.cfs.throttled_time	long	seconds	long	CPU throttled time	the total amount of time that processes have been throttled in the container cgroup
containerengine.cpus.shares	long		long	CPU shares	represents the weight of the cgroup that translates into the amount of CPU it is expected to get. Upon cgroup creation each group gets assigned a default of 1024
containerengine.cpus.quota	long	microseconds	long	CPU quota	enforces a hard limit to the CPU time allocated to processes
containerengine.cpus.period	long	microseconds	long	CPU period	is the time window expressed in microseconds that represents the period for which processes are allowed to run under specific quota
containerengine.network.bytes	long	bytes	long	network received	received amount of bytes on the network interface
containerengine.network.rx.packets	long	packets	long	network packets received	received amount of packets on the network interface

Name	Type	Unit	Numeric Type	Label	Description
container_network.rx.errors	network	long		network rx errors	received amount of errors on the network interface
container_network.rx.dropped	network	long		network rx dropped	amount of dropped inbound packets on the network interface
container_network.tx	network	long		network transmitted	transmitted amount of bytes on the network interface
container_network.bytes	network	long		network received	transmitted amount of bytes on the network interface
container_network.tx.packets	network	long		network transmitted packets	transmitted amount of packets on the network interface
container_network.tx.errors	network	long		network tx errors	transmitted amount of errors on the network interface
container_network.tx.dropped	network	long		network tx dropped	amount of dropped outbound packets on the network interface

## More about Docker Monitoring

- Docker Container Monitoring and Management Challenges
- Docker Container Performance Metrics
- Docker Container Monitoring Open Source Tools
- Docker Container Monitoring with Sematext