

Title: Spark Monitoring Integration Monitor your Hadoop ecosystem with Sema-text Infrastructure and Application Monitoring SaaS, namely HDFS, MapRe-duce, YARN, and Spark. Visualize Apache Spark Manager/Master, Worker, Executor, and Driver componenets, and other key metrics reports. Create dash-boards and alerts, and send notifications to your devops team through various messaging integrations

Integration

- Instructions: <https://apps.sematext.com/ui/howto/Spark/overview>

Metrics

Metric Name Key <i>(Type) (Unit)</i>	Description
driver disk used spark.driver.disk.used <i>(long gauge) (bytes)</i>	
driver memory max spark.driver.memory.total <i>(long gauge) (bytes)</i>	
driver memory used spark.driver.memory.used <i>(long gauge) (bytes)</i>	
driver memory free spark.driver.memory.free <i>(long gauge) (bytes)</i>	
active driver job spark.jobs.active <i>(long gauge)</i>	
all driver job spark.jobs <i>(long gauge)</i>	
failed driver stages spark.stages.failed <i>(long gauge)</i>	
running driver stages spark.stages.running <i>(long gauge)</i>	
waiting driver stages spark.stages.waiting <i>(long gauge)</i>	
executor large file read spark.executor.file.reads.large <i>(long gauge)</i>	
executor file reads bytess spark.executor.file.reads.bytes <i>(long gauge) (bytes)</i>	
executor file read spark.executor.file.reads <i>(long gauge)</i>	
executor file writes bytess spark.executor.file.writes.bytes <i>(long gauge) (bytes)</i>	
executor file writess spark.executor.file.writes <i>(long gauge)</i>	
executor large hdfs read spark.executor.hdfs.reads.large <i>(long gauge)</i>	
executor hdfs reads bytess spark.executor.hdfs.reads.bytes <i>(long gauge) (bytes)</i>	
executor hdfs read spark.executor.hdfs.reads <i>(long gauge)</i>	
executor hdfs writes bytess spark.executor.hdfs.writes.bytes <i>(long gauge) (bytes)</i>	
executor hdfs writess spark.executor.hdfs.writes <i>(long gauge)</i>	
executor active task spark.executor.tasks.active <i>(long gauge)</i>	
executor completed task spark.executor.tasks.completed <i>(long gauge)</i>	

Metric Name Key (Type) (Unit)	Description
executor used pool sizespark.executor.used.pool.size (long gauge) (bytes)	
executor max pool sizespark.executor.max.pool.size (long gauge) (bytes)	
spark appsspark.apps (long gauge)	
waiting appsspark.apps.waiting (long gauge)	
spark workerspark.workers (long gauge)	
live spark workerspark.workers.alive (long gauge)	
spark receiverspark.receivers (long gauge)	
total completed batchesspark.batches.completed (long gauge)	
unprocessed batchesspark.batches.unprocessed (long gauge)	
waiting batchesspark.batches.waiting (long gauge)	
running batchesspark.batches.running (long gauge)	
retained completed batchesspark.batches.completed.retained (long gauge)	
total received recordsspark.records.received (long gauge)	
total processed recordsspark.records.processed (long gauge)	
last completed batch scheduling	
timespark.batches.completed.scheduling.delay (long gauge) (ms)	
last completed batch processing	
timespark.batches.completed.processing.delay (long gauge) (ms)	
last received batch scheduling	
timespark.batches.received.scheduling.delay (long gauge) (ms)	
last received batch processing	
timespark.batches.received.processing.delay (long gauge) (ms)	
last received batch recordsspark.batches.received.records (long gauge)	
worker executorsspark.executors (long gauge)	
worker cores usedspark.cores.used (long gauge)	
worker cores freespark.cores.free (long gauge)	
worker used memoryspark.worker.memory.used (long gauge) (bytes)	
worker free memoryspark.worker.memory.free (long gauge) (bytes)	
uncompleted appsspark.applications.uncompleted (long gauge)	