

title: Log Management FAQ description: Cloud based SaaS / On-premises logging as a service platform FAQ on log shipping, log events fields and count, data retention and daily volumes, Kibana and more

General

How many logs can I store?

There is no limit to how many logs you can store in your logging management app. That said, you should try to pick the ideal Daily Volume when picking your plan, as that is going to give you the most optimal \$/GB price.

If you pick too small Daily Volume and often go over the selected volume your \$/GB price will be higher than if you picked the next higher Daily Volume.

How many fields can my log events have?

A single Logs App allows up to 1000 fields. Most of the time you will want to pay attention to how you structure your logs in order to get the most value out of them, and that typically includes having a manageable number of fields in your logs.

We are getting “Too many fields present in the index”. What does that mean and how can we fix it?

This means you have more than the allowed number of distinct fields in your Logs App. To manually remove some fields use the Field Editor, but if you continue shipping logs with fields you’ve removed via Field Editor those fields will be recreated.

Some reasons why this may happen and how to address it: * You have a high number of different types of log events with different sets of fields, maybe from different sources. You may want to consider using multiple Logs Apps in such cases. You can ship logs from different sources to different Logs Apps, ship logs to different Logs Apps based on the matching patterns), or use grep plugin to drop logs that match a certain pattern. * You have many different deeply nested fields (e.g. foo.bar.bar, foo.bar.bam, etc.), perhaps because you log “objects” with attributes. Often times you don’t want all attributes to turn into individual fields. If you are using Logagent you can use REMOVE_FIELDS or remove-fields output filter to remove fields.

What is @timestamp field?

@timestamp is a date field that is used to filter logs by a selected time. If a log event does not have such a field defined, we will add it and set it to the time when we first received the log event.

How can I check how many GB/day I am shipping to your centralized logging management solution?

From the application, click the **App Actions** button and select **Usage**. There's also **Settings** button next to any of your Logs apps in the **Logs > All Logs Apps** section in our centralized logging and monitoring management solution.

On click, **App Settings > Usage** page will present both Total Volume in Gb and Total Count of indexed logs for that particular Logs App.

How can I check the number of logs I currently have in a Logs App?

Total Count of indexed logs is displayed in **App Settings > Usage** page.

The number of logs will be also displayed in the right corner of the Log Counts panel just below the app's search input. It will change as you perform search queries, change time and apply filters in your Logs App.

You can also do that from Kibana by searching for all your logs without adding any time filters. The number of hits represents the number of all your logs.

How long are my logs stored?

How long your logs are stored depends on what Data Retention you selected when you picked your Logging Management App plan. Each such app can have a different Daily Retention, Daily Volume, and plan. You can change retention, volume, and plan at any time.

Click the **App Actions** button and select **App Settings**. App Info page will open where you can see and change Daily Retention, Daily Volume, and Plan for that particular Logs App.

Can I send old logs?

Yes, you can send all logs. However, keep in mind your Logs App's Data Retention settings. For example, if your retention is set to 7 days and you send logs that are 10 days old, we will give you a warning saying that you are outside of your valid retention time range.

How can I send old logs?

Use Logagent. With Logagent you can then do: `cat /path/to/old/log | logagent -t YOUR_LOGSENE_APP_TOKEN`

What is data retention and how does it work?

Data retention mechanism is responsible for removing old log events. If you have retention set to 7 days, all log events that are received more than 7 days ago will be removed. For paid plans we remove whole 24-hour periods once a day. For free plans we remove old data every 30 minutes.

Is data retention based on index/import/upload/receive time or the actual log event @timestamp?

Data retention is based on log events' receive timestamp.

Are logs shipped to Logs App ever rejected?

When you pick a Daily Volume we automatically set another logging system property called *Max* Daily Volume. By default we set it to 2x the selected Daily Volume. When this limit is hit, our centralized logging management solution will stop accepting new logs for that Logs App until the next day (00:00 UTC). This Max Daily Volume is adjustable, so you can set it to whichever value you'd like.

This Max Daily Volume limit prevents run-away logging from causing over-billing. Logs App will send email notifications before this limit is reached.

What is the maximum log size Logs App will accept?

That depends on the plan you selected for your application. If your application is using the *Basic* plan, our logging management platform will accept logs up to 64 KB in size. If your application is using the *Standard* plan, it will allow logs up to 128 KB. If your application is using the *Pro* plan, it will accept logs up to 256 KB in size. If your logs are larger than the mentioned limits, please consider using the *Enterprise* plan.

My logs have special structure. Can Logs App handle that?

Yes, if the default log index fields (also known as index mapping) don't fit your needs you can create completely custom index mapping. See Custom Mapping Template How-To.

I have multiple log sources - should I send them all to the same Logs App?

Sending logs from multiple log sources to the same Logs App is not a problem at all. However, if you have multiple different log structures (different "fields" in the logs), see the next FAQ entry.

I have multiple different log structures, each with a different set of fields. How should I handle that?

If you have N different log structures, the best way to handle that is by creating N Logs Apps, each with its own index mapping. For example, you may have web server logs, your system logs in /var/log/messages, and your custom application logs. Each of these 3 types of logs has a different structure.

The web server logs probably use Apache Common Log format, the logs in /var/log/messages have syslog structure, and your own application's logs can

be in any format your application happens to use. To handle all 3 log formats elegantly simply create 3 separate Logs Apps and use a different format for each of them.

If you want to be able to search all logs together, even when they have different fields, you can do that - just ship them all to the same Logs App. The thing to watch out for are fields that have the same name, but different data type. For example, if one log source has a purely numeric field “size” and another log source also has a field “size”, but can have non-numeric values, this will cause issues.

Thus, either keep logs of different structures in separate logging app or make sure there are no conflicting fields.

My data structure changed, can I reindex my data automatically?

Yes, when you change your index structure using Field Editor you will be given the opportunity to reindex your data.

Please keep in mind that the speed of reindexing depends on how much data needs to be reindexed and that reindexing is treated like original ingestion, which means it counts towards your daily volume and could result in you hitting your Max Daily Volume Limit.

Is automatic reindex free?

Using Sematext’s Logs App reindex functionality is available for all plans and the normal plan charges apply. Because reindexing counts just like the original data ingestion, before reindexing the estimated cost of reindexing is shown in the UI.

If the amount of data ingested for the day, plus data to be reindexed is smaller than your Max Daily Log Volume Limit then no additional charges will be applied.

Log Shipping

Which log shippers, logging libraries, and platform integrations are supported?

You can use numerous log shippers, logging libraries, and platform integrations.

How can I import logs from my existing Elasticsearch / ELK / Elastic stack?

You can use Logstash to copy logs from your own ELK / Elastic Stack to our Logs App as described at Elastic Stack Import-Export with Logstash & Logsene.

How can I ship logs from Android and iOS apps?

For shipping logs from Android apps use <https://github.com/sematext/sematext-logsene-android> and for shipping logs from iOS apps use <https://github.com/sematext/sematext-logsene-ios>.

How can I ship AWS CloudTrail logs, AWS VPC logs, and other logs from S3?

Use the <https://github.com/sematext/sematext-logsene-aws-lambda-cloudwatch> AWS Lambda function. See Forwarding CloudTrail or Other Logs from AWS S3 to Logsene and Sending AWS Cloudwatch Logs to Logsene for detailed description of how to do this.

Why is it that I can forward logs with rsyslog via TCP or UDP, but not via RELP?

You probably need to install the omrelp module first. If you have rsyslog installed in your system, there's probably a package that you can install:

For RedHat and CentOS:

```
% yum install rsyslog-relp
```

For Debian and Ubuntu:

```
% apt-get install rsyslog-relp
```

For SUSE and OpenSUSE:

```
% zypper install rsyslog-module-relp
```

If you compiled rsyslog manually, you need to supply `--enable-relp` to the configure script. You can check if rsyslog starts properly by starting it in foreground:

```
% rsyslogd -n
```

You will get an error if the module is not found or you have any other problem. For debug mode, add `-d` as well.

How do I ship exception stack traces and other multi-line logs with Logstash without breaking them into multiple log events?

Please see Handling Stack Traces with Logstash.

Plans & Prices

How much does your centralized logging management solution cost?

Check the Logs App Plans & Prices.

What happens with my logs when I change from the free to a paid plan or vice versa?

You don't have to worry about your data when you switch from the free plan to a paid one. Your logs will remain safe and you can still be able to search them. Your data will remain accessible based on your data retention settings. However, when you move from a paid plan to a free plan, free plan restrictions will be applied to your data, which including data retention, too.

What happens if I go over the selected daily log volume and how much does that cost?

The GB/day volume is **not** a limit. That number controls only *cost per GB*. When you go over the GB/day number you had selected we do not reject your data (for protection, use Max Daily Volume). Instead, we keep track of how much data over the selected daily volume was shipped and add that to the monthly cost using the \$/GB rate for the selected daily volume.

The higher daily volume you select, the lower the per GB price. For example, if you picked 1 GB/day and you typically don't go over it, or go over very little, then it pays to keep that 1 GB/day selection. However, if you ship closer to 5 GB/day, then you should select 5 GB/day volume, because that will give you a lower \$/GB rate than if you stay with 1 GB/day selection.

How much data can I ship to Sematext using the Basic plan?

Each day, a total of 500 MB/day can be shipped to all Sematext Logs Apps together. To ship more data use one of the bigger plans.

Troubleshooting

Why am I not seeing my logs?

The most common reason is that logs' timestamps are in the future, and this most often happens when timestamps don't specify the timezone. For example, if you are east of UTC (e.g. in +0100 – 1 hour ahead of UTC) and you send logs with timestamps in your local time, but without specifying your local timezone, like this: 2016-01-30T14:00:00.000Z – this timestamp actually represents 14:00 in UTC. However, if it 14:00 in your timezone, which is 1 hour ahead of UTC, that means that 14:00 UTC is actually 1 hours in the future.

Currently, our Logs App does not display logs that have timestamps in the future, but it does accept them, and shows them when those timestamps stop being in the future.

Another possibility why you are not seeing your logs is that their structure clashes with the structure of your existing logs in the same Logs App. For example, if you have sent logs with field “priority” as integer field then Sematext's Logging Platform will expect all future values of “priority” field to be integers.

Sending a string value in a “priority” field will fail. Note that you *can* change the underlying field “mapping” at any time - see elsewhere in this FAQ.

If none of the above turns out to be a problem in your case, please let us know via support@sematext.com or via our live chat!

How to check if there are any network connectivity issues?

Things to check to ensure network connectivity is OK:

1. Try to connect to `logsene-receiver.sematext.com` / `logsene-receiver.eu.sematext.com` (port 80 or 443) with the following command:

```
nc -zv -w 20 logsene-receiver.sematext.com 443
```

or if you are using syslog, connect to `logsene-syslog-receiver.sematext.com` / `logsene-syslog-receiver.eu.sematext.com` (port 514 for TCP and UDP, port 20514 for RELP, port 10514 for TLS):

```
nc -zv -w 20 logsene-syslog-receiver.sematext.com 10514
```

The output should show something like: Connection to `logsene-receiver.sematext.com` 443 port [tcp/https] succeeded!

In case you see some other result:

- if your server requires proxy to access the Internet, you’ll have to adjust your log shipper config to use proxy settings
- if firewall is used to protect your server, it may be blocking outbound traffic from it. Depending on which receiver and protocol you are using, please ensure with your IT/network admins that right ports are open for outbound traffic
- check your DNS (see below)

2. Check if your DNS has correct entries for Logs App Receiver:

```
nslookup logsene-receiver.sematext.com
nslookup logsene-receiver.eu.sematext.com
```

The output of these commands should look like this, although the IP addresses and names may be somewhat different, as they change periodically:

```
Server:      127.0.1.1
Address:     127.0.1.1#53
```

Non-authoritative answer:

```
logsene-receiver.sematext.com    canonical name = VPC-Logsene-Token-Receiver-LB-188765535.us-east-1.elb.amazonaws.com
Name:      VPC-Logsene-Token-Receiver-LB-188765535.us-east-1.elb.amazonaws.com
Address:   52.54.165.65
Name:      VPC-Logsene-Token-Receiver-LB-188765535.us-east-1.elb.amazonaws.com
Address:   34.196.180.172
```

Or, if you are using syslog, try:

```
nslookup logsene-syslog-receiver.sematext.com
nslookup logsene-syslog-receiver.eu.sematext.com
```

The output should be similar to:

```
Server:      127.0.1.1
Address:     127.0.1.1#53
```

Non-authoritative answer:

```
logsene-syslog-receiver.sematext.com    canonical name = ec2-52-44-248-43.compute-1.amazonaws.com
Name:   ec2-52-44-248-43.compute-1.amazonaws.com
Address: 52.44.248.43
```

How to avoid ‘failed events’?

Sematext Cloud is a schema-less storage and uses Elasticsearch as log storage. When your App receives new log lines, containing unknown fields, those fields are automatically created with the data type of the field value (string, number, object). So you don’t need to specify a schema (called ‘mapping’ in Elasticsearch) upfront. When you ship logs (Elasticsearch document) with the same field name and different types (string, number, object) an error called `mapper_parsing_exception` happens in Elasticsearch. This means the document can’t be indexed because of a mismatch in the data schema. Sematext Cloud catches this error and produces a **failed event** entry in your Logs App. The **failed event** entry contains the error message and the original document as JSON string in the field `logsene_orig_log`.

You can resolve the problem with the `mapper_parsing_exception` with the following procedure: 1) Check the `logsene_error` message, it contains details which field caused the conflict 2) Make sure that the fields use only one type (not a mix of string, number, object ...)

You could do three different things to avoid schema conflicts: 1) Adjust your log output to match the schema. E.g. rename the fields in your logs. 2) Ship logs with different schema to different Sematext Cloud Log Apps.

Example:

Service A produces a log with a field “detail” with a text message.

```
{"detail": "Some error details here ..."} 
```

Service B produces a log with a field “detail” with a JSON object

```
{"detail": {"code": -1, "error": "Some Error message here", "module": "mycode"}}
```

What could be done in this situation? There are basically two options:

- 1) You could rename the field in the log output of in “Service B” e.g. `detail` -> `detail_object`.


```
{"detail_object": {"code": -1, "error": "Some Error message here", "module": "myapp"}
```

The renaming could take place in your apps logging code, or in a log shipper configuration. Some log shippers have support to rename fields before logs get shipped.

- 3) You could ship the logs of “Service B” to a separate Logs App (different App token). Create a new Logs App in Sematext UI and configure your log shipper to ship logs from “Service B” to the new Logs App.

Option #2 is probably the simplest way to resolve the issue.

In case you created only one time logs with a wrong structure, you could use the Field Editor to change the schmemma/mapping. Please refer to related FAQ section: I have multiple different log structures, each with a different set of fields. How should I handle that

Security

Can I ship data to your centralized logging management solution using a secure and encrypted transport like TLS/SSL or HTTPS?

Yes, all data can be shipped both over encrypted channels like TLS/SSL and HTTPS, or via their unencrypted counter-parts. See <https://apps.sematext.com/ui/logs> for more information about how to configure log shipping.

How secure is my data?

Our Logging Management Platform runs on Amazon AWS, whose infrastructure follows IT security best practices and adheres to a number of compliance standards such as SOC 1/SSAE, SOC 2, SOC 3,16/ISAE 3402 (formerly SAS 70 Type II), FISMA, DIACAP, FedRAMP, and HIPAA. For more information about AWS security and compliance certifications see: <http://aws.amazon.com/security> and <https://aws.amazon.com/compliance/>.

Your data is protected from unauthorized reads. If you want to read data using scripts you can authorize your API calls using `apiKey`. You can read more about this [here](#) and [here](#).

Where is my data stored?

Logs App runs and stores data in Amazon AWS in the US and the EU. You can choose if you want your account and your data in the US or in the EU.

How long is my data stored?

This depends on the data retention you selected for each of your logging apps. Backups are stored in S3 and kept for 90 days.

Kibana

Can I use my own Kibana?

Yes.

How can I get a nice map of the world in Kibana?

Ensure you have a country field in your logs. If you only have IP you can use Logstash geoip filter. For example, a configuration like this:

```
input {
  stdin {
    type => "human"
  }
}
filter {
  geoip {
    source => "message"
  }
}
output {
  elasticsearch_http {
    host => "logsene-receiver.sematext.com"
    port => 80
    index => "YOUR-APPLICATION-TOKEN-GOES-HERE"
  }
}
```

If you then start Logstash and type in an IP, you'll see in Kibana an object called **geoip**, which contains lots of information, including country codes. You can then use **geoip.country_code2** as your field in Kibana.

Why are new fields not visible in Kibana and how do I fix it?

Kibana doesn't update field lists automatically. This problem is not specific to our Logs App, but is a general Kibana issue. If you add new fields to an index you must refresh the fields in Kibana. New fields will appear and will be ready to use in your visualizations and searches.

How do I create a dashboard in Kibana?

Kibana requires the following high level steps to create a dashboard:

1. Create a visualization - click to 'Visualize' and save it when finished. See also: <https://www.elastic.co/guide/en/kibana/current/visualize.html>
2. Create a "Search" in the Discover view and save it. Once you do that you can add it as table to your Dashboard later (See screenshot "Log View"). See also: <https://www.elastic.co/guide/en/kibana/current/discover.html>

3. Go to the Dashboard view and add the saved visualizations to it. See also: <https://www.elastic.co/guide/en/kibana/current/dashboard.html>
4. Save the Dashboard

Why can't I draw charts (dashboards) using a given field?

By default all string values sent to Logs App are analyzed. For each analyzed fields we also create `.raw` field which is not analyzed. If you want to use string field for drawing charts you should either use `.raw` field or change index templates to use not `_analyzed` type (<https://sematext.com/blog/custom-elasticsearch-index-templates-in-logsene/>). Let's consider an example. We have a field called 'title'. Our logging management solution uses this field as analyzed one so it is not possible to use it in dashboards. It also creates automatically a field called 'title.raw' and this one can be used for charts.

Timestamps

Why does my syslog timestamp have a different time zone?

The original syslog specification (RFC 3164) has a date field that doesn't contain any time zone information. Fortunately, there's a new specification (RFC 5424) that uses ISO 8601 timestamp. That's why we recommend using the new format for syslog daemons that support it, such as rsyslog or syslog-ng. If you use an older syslogd implementation, or you simply prefer to use the older format, you have two options to have a correct timestamp:

- set the system time on your server to UTC, or
- change the timestamp for your logs to ISO 8601. For example, rsyslog does it with the default template `RSYSLOG_ForwardFormat`. For instance, to forward all logs via TCP, with the original specification, but with the ISO timestamp, you'd add the following line to your config:
`. @@logsene-syslog-receiver.sematext.com:514;RSYSLOG_ForwardFormat`

What are the supported timestamp formats?

If you send your logs via the Elasticsearch API, you can use:

- **ISO8601**. For example, **2001-06-08T08:00:01.123Z** or **2001-06-08T08:00:01+03:00**. A timezone must be there after the time (use **Z** for UTC), so *2001-06-08T08:00:01.123* won't work. You can also send timestamps with just the date (**2001-06-08**), and we'll count them as occurring at 00:00
- **UNIX time in milliseconds**, for example **1487854172123**. UTC is assumed here, so we recommend setting the system time to UTC if you go this route

If you send your logs via syslog, you can use **ISO8601**, too (recommended, see the previous question) or the legacy **RFC-3164 timestamp** (like **Oct 11 22:14:15**). The legacy format assumes UTC and assumes the current year. It doesn't have sub-second precision, either.

See Supported Date Formats for more examples of supported date and time formats.

What happens if I try to send data with @timestamp older than my retention time?

You can send old logs because our data retention works by checking the received timestamp, not @timestamp. For example, if your retention is set to 7 days and you send logs that are 10 days old (i.e. they contain a timestamp from 10 days ago), they will get accepted and processed, but when you set a time range to select those logs, you might get warning that you selected time range that is outside of your retention.

What happens if I try to send data with @timestamp that is in the future?

You can send logs with timestamps in the future, but they will not be shown in Logs App until their timestamps stops being in the future. For example, if you send logs that are “2 hours ahead of real time” they will be accepted, but will not be displayed until 2 hours later.

Alerts

What are Logs App Alerts?

Logs App alerts are based on saved searches that trigger alert notifications when these saved searches meet a certain threshold condition in your logs - e.g. less than N matches in 5 minutes, more than M matches in 3 minutes, etc., or when such saved searches detect *sudden changes* in the number/volume of matching logs - i.e., anomalies.

How to create Alerts in Logs App?

Alerts are added by clicking Save Query icon located next to the Logs App input field. Enter query info and select option to enable alerts. Alerts panel will expand and allow you to customize alert specific to that particular search query.

How to view Alerts?

Click on Saved Queries Icon located next to the app's search input field.

Saved Queries page will open and your queries and related alerts will be displayed here

What is the difference between threshold-based Alerts and Anomaly Detection (aka Algoterts)?

If you have a clear idea about how many logs should be matching a given Alert Query, then simply use threshold-based Alerts. In other words, if you know that you always have some ERROR-level log events, you may want to use threshold-based alerts for that.

For example, if your logs typically have < 100 ERROR-level messages per minute you may want to create an Alert Query that matches ERROR log events and notifies you when there are more than 100 such matches in 1 minute.

If you do not have a sense of how many matches a given Alert Query matches on a regular bases, but you want to watch out for sudden changes in volume, whether dips or spikes, use Algoterts (Anomaly Detection-based Alerts).

An extreme example that you may actually like is an Alert Query that matches all logs - "*" - a single wildcard character for a query. If you use that for an Algotert then Logs App will notify you when the overall volume of your logs suddenly changes, which may be a signal that you should look for what suddenly increased logging (e.g. maybe something started logging lots of errors, exceptions, timeouts, or some such), or why logging volume suddenly dropped (e.g. maybe some of your server(s) or app(s) stopped working and sending logs).

Can I send Alerts to Slack, PagerDuty, Nagios, or other WebHooks and ChatOps services?

Yes. Our centralized logging management solution lets one configure arbitrary WebHooks to call with Alert event information when Alert events are triggered. See alerts FAQ for more info.

```
<div class="header">
  <a href="#" id="close">
    <div class="box box3">
      <svg viewBox="0 0 40 40">
        <path class="close-x" d="M 10,10 L 30,30 M 30,10 L 10,30"></path>
      </svg>
    </div>
  </a>
  <h2>Logs App - Usage Menu Link</h2>
</div>
<div class="copy">
  
```